

Scalable Multi-Tenant Architecture for AI-Enhanced Nonprofit Moves Management

Executive Summary of the Architectural Paradigm

The transition from a service-based consulting model to a productized, multi-tenant Software-as-a-Service (SaaS) architecture necessitates a fundamental redesign of how automation and artificial intelligence interact with client data. In the current operational model, the consultant duplicates an Airtable template for each new nonprofit client and establishes discrete automation pathways—such as the existing Zapier proof-of-concept—for every individual instance. While this decentralized approach functions for a limited client roster, it rapidly introduces insurmountable maintenance debt. Pushing a single structural update, modifying an AI prompt, or altering a workflow across fifteen or more fragmented client bases requires exponential labor and introduces severe version control vulnerabilities.¹

To build a sellable product that can be deployed efficiently across multiple clients without requiring separate configurations for each instance, the architecture must decouple the cognitive processing and automation logic from the data storage layer. By establishing a centralized hub-and-spoke infrastructure, a single "master" automation orchestrator and a single set of AI agents can dynamically serve an infinite number of distributed client Airtable bases.¹

This comprehensive implementation plan details the technical execution required to integrate a Donor Research Agent, a Meeting Preparation Agent, and a Move Suggestion Agent into the existing Airtable moves management template. The design leverages advanced webhook routing, dynamic API credentialing, and Retrieval-Augmented Generation (RAG) to achieve enterprise-grade scalability, ensuring that onboarding a new client requires minimal setup time while maintaining a single, unified backend infrastructure.

Platform Evaluation for the Hub-and-Spoke Infrastructure

Achieving true multi-tenancy without writing custom middleware requires the strategic pairing of an automation orchestrator (the central nervous system) and a dedicated AI agent builder (the cognitive engine).

The Orchestration Layer: Make.com vs. Zapier

The consultant has successfully utilized Zapier for an initial proof of concept; however, Zapier

presents critical structural and economic limitations for a sellable, multi-tenant product. Zapier’s pricing model charges per completed task, meaning high-volume workflows rapidly erode profit margins as the client base scales.² Furthermore, Zapier's predominantly linear workflow structure struggles with the dynamic routing required to process requests from dozens of different Airtable bases through a single scenario.

Make.com represents the optimal orchestration platform for this architecture. Make.com offers highly cost-effective scaling, advanced visual workflows, branching logic, and the ability to dynamically route data using complex JSON payloads and API calls.² Crucially, Make.com supports "Dynamic Connections" on its enterprise tiers, and even on lower tiers, it allows for raw HTTP modules to dynamically target different Airtable bases based on incoming webhook parameters.³ This functionality is the cornerstone of the single-backend architecture.

The Cognitive Layer: Relevance AI

While direct integrations with the OpenAI or Anthropic APIs are possible, utilizing a dedicated no-code AI agent builder significantly reduces the technical overhead of managing tool calling, prompt versioning, and long-term memory. Relevance AI stands out as the premier platform for multi-tenant SaaS deployments. Unlike rigid, stateless chat interfaces or basic prompt-chaining tools, Relevance AI allows for the creation of autonomous agents equipped with custom tools (such as web search and API querying) and custom Knowledge tables (vector databases).⁶

Economically, Relevance AI utilizes a pricing model highly conducive to white-labeling and multi-tenant scaling. The platform separates "Action Credits" (the cost of running the agent workflow) from "Vendor Credits" (the actual cost of the Large Language Model computation). Relevance AI passes the LLM costs through without markup, and critically, permits product creators on paid tiers to bring their own API keys (e.g., from OpenAI) to bypass vendor credits entirely.⁹ This structure ensures that as client usage increases, the core computational costs remain at baseline wholesale rates, preserving the consultant's profit margins.

Orchestration Platform	Pricing Structure	Dynamic Routing Capabilities	Verdict
Zapier	High; per-task billing	Limited; favors linear, single-tenant flows ²	Unsuitable for scaling multi-tenant SaaS.
Make.com	Low; high-volume operation tiers	Advanced; JSON mapping, dynamic HTTP ²	Optimal selection for the hub orchestrator.

Cognitive Platform	Primary Strengths	Cost Structure	Verdict
Custom Code (Python/Node)	Infinite flexibility	High development and maintenance overhead	Unnecessary technical burden.
Relevance AI	Visual tool builder, native vector DBs ¹⁰	Flat subscription + wholesale LLM costs ⁹	Optimal selection for the cognitive engine.

Architectural Mechanics of the Hub-and-Spoke Model

To maintain a single manageable backend, the system utilizes a stateless processing mechanism. The "spokes" are the individual, isolated client Airtable bases containing donor data. The "hub" is the single Make.com scenario and Relevance AI setup.

1. **The Dynamic Trigger:** Instead of relying on native Airtable automations that must be manually configured inside every individual client base, each client base utilizes an Airtable Button field.¹²
2. **Payload Delivery:** The Button field uses a formula to trigger a single, centralized Make.com Webhook URL.¹⁴ The formula dynamically appends the specific client's Base ID, the Table ID, and the Record ID to the webhook payload.¹⁵
3. **Centralized Processing:** Make.com receives the webhook, parses the URL parameters to identify which client base triggered the request, and passes the relevant donor data to Relevance AI.
4. **Dynamic Write-Back:** Once Relevance AI completes its analysis, Make.com utilizes an HTTP module to push the enriched data back to the exact client base and record that initiated the request.¹⁷ Native Make.com Airtable update modules cannot accept dynamically mapped Base IDs; therefore, the raw HTTP module is a mandatory architectural requirement.³

This infrastructure guarantees that when the consultant improves an AI prompt or adds a new research data source, the update is instantly deployed to all clients simultaneously, fulfilling the requirement for maximum deployment efficiency and zero per-client maintenance overhead.

Feature 1: The Donor Research Agent

The Donor Research Agent acts as an autonomous prospect researcher, designed to ingest minimal donor details (e.g., name, employer, location) and return a comprehensive briefing containing wealth indicators, philanthropic history, and professional background.

API Data Source Evaluation and Selection

To provide high-value, actionable insights, the agent must query both structured philanthropic databases and unstructured web data.

Structured Philanthropic Data: DonorSearch vs. iWave For nonprofit prospect research, DonorSearch and iWave operate as the industry standards. iWave offers holistic donor profiles combining wealth and philanthropic behavior; however, it caters heavily to enterprise organizations with opaque pricing models and lacks transparent, accessible tiers for low-volume API usage.¹⁸ DonorSearch is highly recommended for this application. It boasts a 90% accuracy rate, focuses deeply on charitable giving history (the strongest predictor of future giving), and offers a robust, well-documented API.¹⁸ Furthermore, DonorSearch API pricing starts at an accessible \$600/year for 750 calls, making it highly suitable for a productized offering where API costs must be modeled into the subscription fee.²²

Unstructured Web Data: Tavily vs. Perplexity

To gather recent news, social media profiles, and employment updates, an AI-optimized search API is required. Traditional SERP APIs (like Google Custom Search) return raw HTML that overwhelms LLM context windows.

Tavily is specifically engineered for LLMs and AI agents, returning clean, extracted content (snippets, markdown, and citations) ready for immediate ingestion.²³ Tavily costs approximately \$8.00 per 1,000 requests and includes integrated extraction, making it highly efficient for real-time donor research.²⁴ While Perplexity’s Sonar API is faster and occasionally cheaper (\$5 per 1,000 requests), it lacks built-in advanced extraction tooling, requiring additional processing layers to format the output for complex downstream agent tasks.²³ Tavily's structured extraction reduces the need for complex parsing logic in Make.com, making it the superior choice.

Data Source	Primary Intelligence Vector	Pricing Model	Architectural Benefit
DonorSearch API	Philanthropic history, wealth screening	Starts ~\$600/year (750 calls) ²²	High accuracy (90%); API built for automated prospect research. ²⁰
Tavily API	Real-time web search, news,	~\$8.00 per 1,000 requests ²⁴	AI-ready markdown extraction prevents

	professional data		LLM token bloat. ²³
--	-------------------	--	--------------------------------

Step-by-Step Technical Implementation

Step 1: Configuring the Client Airtable Base (The Spoke) Within the standard Airtable template provided to the client, navigate to the "Donors" table. Create a new formula field named Base_ID. While Airtable does not have a native BASE_ID() formula function, the consultant can mandate that upon deployment, the client pastes their unique Base ID (found in the Airtable URL, starting with app²⁶) into a static text field, or use complex workarounds to extract it.¹⁶ For reliability, a static configuration field populated once during onboarding is the most stable approach.

Next, create a Button field named "Run AI Research." Set the button action to "Open URL".¹³ The formula for this button dynamically constructs the webhook payload:

CONCATENATE("https://hook.us1.make.com/your_unique_webhook_id?", "recordId=", RECORD_ID(), "&baseId=", {Base_ID}, "&donorName=", ENCODE_URL_COMPONENT({Donor Name}), "&employer=", ENCODE_URL_COMPONENT({Employer}), "&city=", ENCODE_URL_COMPONENT({City}))

When a fundraiser clicks this button, a background browser tab briefly opens, seamlessly transmitting the precise location of the record and the baseline search parameters to the central hub.¹⁵ Make.com can be configured to return a 200 OK status with a simple HTML script to auto-close the tab, ensuring a smooth user experience.

Step 2: Constructing the Relevance AI Agent In the Relevance AI dashboard, create a new Custom Agent using the "Agent Builder".²⁸

- Core Instructions:** Define the agent's persona and constraints. *"You are an expert nonprofit prospect researcher. Your objective is to compile a comprehensive summary report on a donor. You will receive a name, location, and employer. You must use the Tavily Search Tool and the DonorSearch Tool. Synthesize the raw data into a structured markdown report avoiding hallucinations. If data is missing, explicitly state 'Data unavailable'."*³⁰
- Tool Integration:** Utilize Relevance AI's API integration features to construct two custom tools accessible to the agent.
 - Web Search Tool:* Connect the Tavily API endpoint via HTTP request. The agent autonomously generates the search query string based on the donor's name and employer.³²
 - Wealth Search Tool:* Connect the DonorSearch API endpoint. Map the donor demographic variables to the API request parameters to fetch structured giving history.³⁴
- Formatting Directives:** Instruct the agent to format its final response into specific sub-headers: Executive Summary, Professional Background, Philanthropic Footprint, Wealth Indicators, and Recent News.

Step 3: The Make.com Orchestration Workflow

1. **Webhook Trigger:** Create a "Custom Webhook" module. This generates the unique URL used in the Airtable button formula.³⁵ Ensure the module is set to automatically determine the data structure based on the incoming query parameters (recordId, baseId, donorName, etc.).
2. **Relevance AI Module:** Add the Relevance AI module configured to "Trigger Agent." Map the dynamic variables from the webhook payload into the agent's input fields.²⁹
3. **Airtable Update Module (Dynamic HTTP):** As previously established, native Make.com modules cannot update records if the Base ID changes dynamically per request.³ The architecture must utilize the "HTTP - Make a Request" module to interact directly with the Airtable REST API.³⁷
 - o **Method:** PATCH¹⁷
 - o **URL:** `https://api.airtable.com/v0/{{1.baseId}}/Donors/{{1.recordId}}`¹⁷
 - o **Headers:**
 - Authorization: Bearer
 - Content-Type: application/json³⁹
 - o **Body Type:** Raw
 - o **Body (JSON):**

```
JSON
{
  "fields": {
    "AI Research Summary": "{{2.agentOutput}}",
    "Research Last Updated": "{{now}}"
  },
  "typecast": true
}
```

The "typecast": true parameter is vital; it instructs the Airtable API to perform best-effort automatic data conversion (e.g., converting a string timestamp into a proper date field), ensuring maximum system stability across varying client configurations.⁴⁰

Feature 2: The Meeting Preparation Agent

While the Donor Research Agent focuses on external data acquisition, the Meeting Preparation Agent excels at internal data synthesis. It combines the newly acquired external research with the historical engagement data already housed within the client's Airtable CRM to brief a fundraiser immediately prior to a major gift solicitation.

Technical Architecture and Context Window Management

The primary technical challenge with this feature is compiling all relevant historical moves (engagement actions, meeting notes, email summaries) linked to a specific donor and feeding them into the LLM's context window without exceeding token limits or introducing hallucination

risks.

Step 1: Advanced Payload Generation in Airtable

When the user clicks the "Generate Briefing" button in the Airtable Donor record, the webhook payload must carry more than just basic demographics. It requires the historical context. Retrieving this via sequential API searches within Make.com would require complex arrays and multiple API calls, increasing latency and cost.

To achieve data compression prior to transmission, the architecture utilizes Airtable's native Rollup fields. In the Donor table, create a Rollup field that targets the linked "Moves" table. Using the ARRAYJOIN() aggregation formula, the system can compile all previous meeting dates, action types, and notes into a single, comma-separated text string localized directly on the donor record.⁴²

The Airtable button's formula for this agent will append this newly concatenated history:

```
"&history=", ENCODE_URL_COMPONENT({Rollup_Moves_History})
```

```
"&research=", ENCODE_URL_COMPONENT({AI Research Summary})
```

Step 2: Relevance AI Prompt Engineering

The Meeting Preparation Agent in Relevance AI does not require external API tools (like Tavily), as all necessary data is pre-aggregated and delivered via the webhook payload. Its strength relies entirely on sophisticated system prompts and the reasoning capabilities of the underlying LLM.

System Prompt Design:

"You are a strategic major gift advisor preparing a nonprofit fundraiser for a critical meeting. You will receive a donor's historical engagement log (past interactions) and an external research summary. Your task is to generate a highly tactical, actionable meeting brief. Your output must strictly adhere to the following markdown structure:

1. **Executive Summary:** A concise, two-sentence overview of the donor's current disposition and capacity.
2. **Strategic Objectives:** Define the primary goal of this specific meeting based on where they are in the cultivation cycle.
3. **Talking Points:** Provide three specific, personalized conversation starters. One must be based on recent news/business, and one must reference a past engagement from the historical log to demonstrate continuity of relationship.
4. **Strategic Recommendations:** Suggest a potential ask amount or the optimal programmatic project to highlight, derived directly from their philanthropic history."

Step 3: Make.com Routing and Write-Back

Following the identical architectural pattern established in Feature 1, Make.com receives the webhook, passes the structured context (Research Summary + Rollup History) to Relevance AI, and awaits the synthesized brief.

Once received, the Make.com HTTP module executes a PATCH request to the dynamically provided `baseId` and `recordId`. The JSON payload updates a Long Text (rich text enabled) field named "Latest Meeting Brief" in the Airtable Donor record.

Ease of Setup and Maintenance Dynamics

Because this agent relies exclusively on advanced prompt engineering and data already localized within the Airtable payload, its maintenance burden is virtually zero. It does not incur external API costs (saving the \$8/1000 Tavily fees), consuming only standard LLM vendor credits via Relevance AI.⁹ This creates a high-value, low-cost feature that significantly bolsters the perceived utility of the productized template without adding fragile moving parts to the backend.

Feature 3: Move Suggestion and Creation Agent

The "Inspire Me" feature represents the most technologically sophisticated component of the proposed architecture. It moves beyond text summarization and enters the realm of prescriptive analytics by evaluating the extensive library of 70+ pre-built move ideas and intelligently matching 3 to 5 appropriate next steps to a specific donor profile.

Vector Databases and Semantic Matching Logic

Traditional relational database queries and if/then logic are fundamentally insufficient for this task. A donor profile containing qualitative notes such as "Passionate about marine conservation, recently sold a B2B SaaS startup, prefers intimate dinners over large galas, highly analytical" cannot be easily matched to a static list of 70 moves using standard categorical filters. The solution requires semantic vector matching—converting both the donor profile and the text of the moves library into high-dimensional mathematical embeddings, and calculating the cosine similarity between them to discover the most conceptually relevant matches.

Implementing the Knowledge Base in Relevance AI Relevance AI provides native, fully managed infrastructure for this exact use case through its "Knowledge" features and "Vectorize and search array" tools.⁸

1. **Centralized Moves Library:** Instead of storing the 70+ moves in every single client's Airtable base (which would require the Make.com scenario to search a different table for every client), the consultant maintains the master moves library entirely inside Relevance

AI as a global Knowledge Table.¹⁰ Each row in this table contains the Move Name, Move Description, Effort Level, and an Ideal Donor Persona description.

2. **Automated Vectorization:** Relevance AI automatically processes this text into vector embeddings.¹¹ By storing the library centrally, the consultant ensures that when they invent a new strategic "move," they only append it to the Relevance AI Knowledge Table once. It instantly becomes available to the AI agents serving all 15+ client instances, perfectly preserving the multi-tenant SaaS philosophy.

The Agentic Retrieval Workflow

1. **Trigger:** The fundraiser clicks the "Inspire Me" button on a donor record. The webhook payload transmits the donor's demographic data, wealth score, and AI Research Summary to Make.com, which forwards it to the Relevance AI agent.
2. **Semantic Search Generation:** Within Relevance AI, the agent utilizes the "Retrieve records from Knowledge" tool.⁴³ The agent analyzes the donor profile and converts it into an optimized search query. For example, it translates the raw profile into the semantic query: "High capacity donor, environmental focus, introverted engagement strategies, entrepreneurial background."
3. **Evaluation and Filtering:** The vector database retrieves the top 10 semantically similar moves from the central Knowledge Table. The LLM evaluates these 10 retrieved options against the donor's historical engagement log (provided in the payload) to filter out any moves that have already been executed recently. It refines the final list to the top 3 to 5 most optimal recommendations, providing a short rationale for each selection.

Automated Move Creation in Airtable

To drastically reduce decision fatigue and administrative friction, the agent does not merely suggest these moves in a text field; it autonomously creates the suggested moves as actual records directly in the client's Airtable "Moves" table, linking them automatically to the specific donor.

This requires a highly sophisticated Make.com routing sequence utilizing iterators and bulk API requests.

1. **Structured Output:** The Relevance AI agent is instructed to output the 3 to 5 recommended moves strictly as a structured JSON array containing keys for moveName, moveDescription, and rationale.
2. **Iteration:** Make.com receives this JSON array and utilizes an "Iterator" module. The Iterator breaks the single array into multiple discrete bundles (one for each suggested move).⁴⁴
3. **Bulk Record Creation:** To optimize API limits, Make.com can either loop individual POST requests, or preferably, use an "Array Aggregator" post-iteration to format a single bulk HTTP POST request to the client's specific Airtable Base ID to create all new records simultaneously.⁴⁰

The Dynamic Bulk POST Request JSON Body:

To create records and link them to the original donor, the JSON payload must utilize Airtable's specific array structure for linked record fields.

JSON

```
{  
  "records":  
  },  
],  
  "typecast": true  
}
```

By mapping the original donorRecordId (captured from the initial webhook payload) into the "Linked Donor" array field, the Airtable API automatically creates the relational link between the newly generated move record and the parent donor record.³⁹ This delivers an unparalleled, seamless user experience: the fundraiser clicks a button, and seconds later, three fully formulated, strategically sound tasks spontaneously populate their pipeline.

Evaluation Criteria and Productization Viability

Transitioning from a localized consulting asset to a scalable, sellable SaaS product requires rigorous evaluation of operational mechanics. The proposed architecture addresses all core evaluation criteria required for commercial viability.

1. Ease of Build and Technical Complexity

The architecture heavily favors no-code visual builders, making it highly accessible. Make.com's visual interface replaces the need for complex Node.js middleware, allowing someone with novice to intermediate technical skills to map variables visually.² Relevance AI completely abstracts the severe complexities of vector databases, embedding generation, and API chaining into a user-friendly drag-and-drop tool builder.⁴⁸ The most technically demanding component is formatting the raw JSON body for the HTTP requests to the Airtable API (detailed above), which requires only a one-time setup and validation.

2. Scalability and Deployment Efficiency

The hub-and-spoke model fundamentally resolves the scaling bottlenecks that plague Airtable consulting models.¹

- **The Legacy Model:** Selling the template to 10 new clients requires duplicating 10 bases, establishing 10 new Zapier accounts, wiring 30 new Zaps, and maintaining discrete API keys across all environments.
- **The SaaS Model:** Selling the template to 100 new clients requires only duplicating the Airtable template. Upon delivery, the client pastes their unique Base ID into a configuration record. The deployment is instantly active.

Because the Make.com webhook URLs remain static and rely entirely on the dynamically passed Base ID to direct traffic, the backend infrastructure scales infinitely without requiring the consultant to construct a single new configuration node.⁴

3. Cost Structure and Margin Protection

Zapier's per-task pricing model heavily penalizes success; the more the client utilizes the AI features, the more the consultant's profit margins shrink.² Make.com offers operations at a fraction of the cost, capable of handling tens of thousands of complex routing operations for a flat, low monthly fee.²

Relevance AI's pricing model is uniquely advantageous for white-label SaaS deployment. By separating action credits from vendor credits, the consultant pays a flat subscription for the platform infrastructure and provides their own OpenAI or Anthropic API keys.⁹ This guarantees that the consultant only pays the wholesale, per-token cost of the LLM computations, avoiding platform markups. External APIs like Tavily (approximately \$8 per 1,000 requests)²⁴ and DonorSearch (\$600 annually)²² represent highly predictable overhead that can be modeled directly into the client's monthly subscription fee (e.g., offering clients a tiered system such as "100 AI Research credits per month").

4. Reliability and Error Handling

Relying on dynamic HTTP modules bypassing native Airtable modules ensures that the system will not catastrophically break when a client renames a table or alters a view, provided the core Table IDs and Field IDs remain consistent.⁴¹ To maximize reliability, the system mandates the use of typecast: true in all Airtable API requests. This instructs Airtable to handle minor data type mismatches gracefully, significantly reducing API rejection rates.⁴⁰

Furthermore, Make.com features robust built-in error routing. If an external API call fails (e.g., the Tavily search endpoint times out), Make.com can catch the error and route a fallback message back to the Airtable record stating: *"AI Research temporarily unavailable due to external server timeout,"* preventing silent failures and reducing support tickets.

5. Security and Data Handling Protocols

Handling sensitive donor data across multiple nonprofit clients requires strict adherence to data isolation and security protocols. In a multi-tenant environment, cross-client data contamination is the highest risk. The hub-and-spoke architecture mitigates this by functioning entirely as a stateless processor.

Relevance AI and Make.com process the data in transit, generate the insight, and immediately push the data back to the client's isolated Airtable instance. The donor data is not pooled or stored in a centralized database long-term, preserving tenant isolation and compliance.¹⁰

For API authentication with Airtable, the consultant must utilize an Airtable Personal Access Token (PAT).⁵⁰ While OAuth is the gold standard for public, user-authorized third-party applications, implementing an OAuth flow introduces massive developmental complexity. For a consultant-managed template distribution model, a PAT with deep scopes—managed by the consultant acting as a centralized Service Account—is highly effective, fully supported by Airtable, and much simpler to implement.⁵¹ It is critical that this PAT is stored securely as a Make.com Connection Variable and is never exposed in client-facing frontend formulas or scripts.⁵²

Technical Limitations and Mitigation Strategies

While the proposed architecture is highly robust, several technical limitations inherent to the selected platforms must be proactively managed to ensure long-term stability.

The Airtable API Rate Limit Constraint

The standard Airtable API enforces a strict rate limit of 5 requests per second per base.⁵³ If a client attempts to trigger the Donor Research Agent on 20 donors simultaneously (e.g., clicking the button rapidly down a list), the Make.com scenario will hit the Airtable API rate limit during the write-back phase, causing the HTTP PATCH or POST requests to fail.

Mitigation: Within the Make.com scenario, the consultant should utilize the "Sleep" module to introduce a 1 to 2-second delay before the final HTTP requests, effectively pacing the data flow. Additionally, implementing Make.com's built-in error handling "Break" directive will automatically intercept a 429 Too Many Requests error, store the failed execution, and retry the payload after a specified delay to comply with the rate limits without losing data.

LLM Hallucinations in Major Gift Contexts

AI agents, regardless of the underlying model, can occasionally hallucinate data. In the context of major gift fundraising, providing a fundraiser with wildly inaccurate wealth data, fabricating a past interaction, or inventing a philanthropic interest can cause severe reputational damage to the nonprofit if acted upon.

Mitigation: The Relevance AI system prompt must include exceptionally strict grounding directives: *"You must only reference facts explicitly provided in the external research payload or the historical log. Under no circumstances should you infer or estimate wealth capacity if hard data is absent. If a specific data point is missing, you must output 'Data unavailable'."*

Furthermore, UI/UX design within the Airtable template is critical. Any field populated by AI should be clearly demarcated using specific naming conventions (e.g., "🤖 AI Suggested Moves" or "🤖 Research Summary"). During onboarding, clients must be explicitly trained to review all AI-generated output as a foundational starting point for human review, not as absolute truth.

Vendor Lock-in and API Deprecation

Relying on external APIs for core product functionality introduces operational risk if those APIs change their pricing structures or deprecate endpoints (e.g., if Tavily alters its markdown extraction format, or DonorSearch adjusts its authentication flow).

Mitigation: The modularity of the architecture provides a strong defense against vendor risk. Because Relevance AI abstracts the LLM and the tools from the orchestration layer, if Tavily becomes economically unviable or technically unstable, the consultant can simply swap it for the Perplexity Sonar API or Exa.ai within the Relevance AI dashboard. The Make.com routing logic and the client's Airtable interface remain completely untouched, ensuring continuity of service without disrupting the frontend client experience.²⁸

Architectural Synthesis

By executing the strategic implementation detailed in this report, the consultant will successfully transition an unscalable, manual Airtable template into an automated, highly lucrative multi-tenant SaaS product.

The integration of Make.com as a dynamic orchestration layer eliminates the compounding maintenance debt associated with managing dozens of individual automation accounts. Simultaneously, utilizing Relevance AI provides the cognitive infrastructure necessary to deploy highly sophisticated, tool-wielding AI agents that deliver immediate, tangible ROI to nonprofit clients.

The resulting product—featuring autonomous donor research via structured APIs, context-aware meeting synthesis, and intelligent, vector-matched move suggestions—represents a paradigm shift in donor moves management. It empowers the consultant to onboard an unlimited number of clients with near-zero technical setup, protecting profit margins, ensuring data security, and solidifying a dominant position as a premier provider of AI-enhanced fundraising technology.

Works cited

1. How are you managing multi-client Airtable bases at scale without duplicating logic, breaking automations, or hitting sync limits? - Reddit, accessed February 22, 2026,
https://www.reddit.com/r/Airtable/comments/1mdxgpe/how_are_you_managing_multiclient_airtable_bases/
2. Airtable Automations vs Zapier vs Make vs n8n: The 2025 Guide - Agility HQ Australia, accessed February 22, 2026,
<https://agilityhq.au/articles/airtable-vs-zapier-make-n8n>
3. Unable to Map Scenario Data to Airtable "Update Record" Module - Make Community, accessed February 22, 2026,
<https://community.make.com/t/unable-to-map-scenario-data-to-airtable-update-record-module/46193>
4. Dynamic connections - Help Center, accessed February 22, 2026,
<https://help.make.com/dynamic-connections>
5. Dynamic Connections | Latest Capability on Make in 2024, accessed February 22, 2026, <https://www.make.com/en/blog/dynamic-connections>
6. No Code AI Agent Builder - FME by Safe Software, accessed February 22, 2026,
<https://fme.safe.com/guides/ai-agent-architecture/no-code-ai-agent-builder/>
7. Relevance AI - Build your AI Workforce - AI for Business, accessed February 22, 2026, <https://relevanceai.com/>
8. Introduction to Agents - Relevance AI Documentation, accessed February 22, 2026, <https://relevanceai.com/docs/agent/introduction>
9. Plans and credits - Relevance AI Documentation, accessed February 22, 2026,
<https://relevanceai.com/docs/get-started/subscriptions/plans>
10. Knowledge - Relevance AI Documentation, accessed February 22, 2026,
<https://relevanceai.com/docs/knowledge/create-knowledge>
11. Vectorize and Search array - Relevance AI Documentation, accessed February 22, 2026, <https://relevanceai.com/docs/tool/tool-steps/vectorize-and-search>
12. Use Airtable Buttons to Run External Automations - YouTube, accessed February 22, 2026, <https://www.youtube.com/watch?v=qGSGnklfeX4>
13. Button Field - Overview - Airtable Support, accessed February 22, 2026,
<https://support.airtable.com/docs/button-field>
14. Trigger script actions from interface - Airtable Community, accessed February 22, 2026,
<https://community.airtable.com/development-apis-11/trigger-script-actions-from-interface-4095>
15. Sending Data to integromat via Webhook using a Button - Airtable Community, accessed February 22, 2026,
<https://community.airtable.com/automations-8/sending-data-to-integromat-via-webhook-using-a-button-25838>
16. There is no way to get the BASE_ID() and TABLE_ID() using a formula? Right?, accessed February 22, 2026,
<https://air.tableforums.com/t/there-is-no-way-to-get-the-base-id-and-table-id-u>

[sing-a-formula-right/1306](#)

17. Update record - Airtable Web API, accessed February 22, 2026, <https://airtable.com/developers/web/api/update-record>
18. iWave vs. DonorSearch: How They Compare : r/Nonprofitprofessional - Reddit, accessed February 22, 2026, https://www.reddit.com/r/Nonprofitprofessional/comments/1ftr62e/iwave_vs_donorsearch_how_they_compare/
19. Top prospect research tools and resources for fundraising - Kindsight, accessed February 22, 2026, <https://kindsight.io/resources/blog/prospect-research-tools/>
20. DonorSearch vs. iWave Kindsight, accessed February 22, 2026, <https://www.donorsearch.net/donorsearch-vs-iwave-kindsight/>
21. Top 7 Donor Prospect Research Software for Secure Nonprofits - Snowball Fundraising, accessed February 22, 2026, <https://snowballfundraising.com/top-donor-prospect-research-software/>
22. DonorSearch vs. EveryAction – Best donor management for nonprofits (2025) | Zeffy, accessed February 22, 2026, <https://www.zeffy.com/compare/donorsearch-vs-everyaction>
23. 5 Tavily Alternatives for Better Pricing, Performance, and Extraction Depth - Firecrawl, accessed February 22, 2026, <https://www.firecrawl.dev/blog/tavily-alternatives>
24. Perplexity Search API vs. Tavily: RAG & Agent Choice 2025 - AlphaCorp AI, accessed February 22, 2026, <https://alphacorp.ai/perplexity-search-api-vs-tavily-the-better-choice-for-rag-and-agents-in-2025/>
25. Best AI-Ready Search APIs for LLMs & Agents in 2025: Serpex.dev vs Tavily vs Bright Data, accessed February 22, 2026, <https://www.serpex.dev/blog/best-ai-ready-search-apis-for-llms-agents-in-2025-serpexdev-vs-tavily-vs-bright-data>
26. Finding Airtable IDs, accessed February 22, 2026, <https://support.airtable.com/docs/finding-airtable-ids>
27. Button Field for Open URL using a Shared Link with Dynamic Filters - Airtable Community, accessed February 22, 2026, <https://community.airtable.com/fid-9/tid-30400>
28. Build and Recruit Autonomous AI Agents - Relevance AI, accessed February 22, 2026, <https://relevanceai.com/agents>
29. Build AI Agents the Easy Way | No-Code Guide with Relevance AI - YouTube, accessed February 22, 2026, <https://www.youtube.com/watch?v=81EzOX9hD6g>
30. Using Airtable AI in fields, accessed February 22, 2026, <https://support.airtable.com/docs/using-airtable-ai-in-fields>
31. How to Build AI Agents Step by Step: Relevance.ai Tutorial - YouTube, accessed February 22, 2026, <https://www.youtube.com/watch?v=KeF8TIRw7jI>
32. Tavily - Relevance AI, accessed February 22, 2026, <https://relevanceai.com/integrations/tavily>
33. Deep Research AI Workflow Using Langgraph + Tavily + Any LLM Provider - Medium, accessed February 22, 2026,

- <https://medium.com/@gaurav219688/deep-research-ai-workflow-using-langgrap-h-tavily-search-any-llm-provider-373ae5aa2cfd>
34. Nonprofits | DonorSearch, accessed February 22, 2026, <https://www.donorsearch.net/ds-for-nonprofits/>
 35. Webhook and Airtable: Automate Workflows with n8n, accessed February 22, 2026, <https://n8n.io/integrations/webhook/and/airtable/>
 36. Make.com/Airtable webhook | Airtable Community, accessed February 22, 2026, <https://community.airtable.com/automations-8/make-com-airtable-webhook-24925>
 37. Using the update record automation action - Airtable Support, accessed February 22, 2026, <https://support.airtable.com/docs/update-record-action>
 38. How to Use HTTP Module in Make.com: My Complete Guide to Custom API Magic - Lets Viz Technologies, accessed February 22, 2026, <https://lets-viz.com/blogs/how-to-use-http-module-in-make-com-my-complete-guide-to-custom-api-magic/>
 39. How to I insert records into link fields with API? - Airtable Community, accessed February 22, 2026, <https://community.airtable.com/development-apis-11/how-to-i-insert-records-into-link-fields-with-api-7126>
 40. Update multiple records - Airtable Web API, accessed February 22, 2026, <https://airtable.com/developers/web/api/update-multiple-records>
 41. Create records - Airtable Web API, accessed February 22, 2026, <https://airtable.com/developers/web/api/create-records>
 42. Pull Airtable linked record data into Make.com - YouTube, accessed February 22, 2026, <https://www.youtube.com/watch?v=7pxeVUdvKQo>
 43. Retrieve Record(s) from Knowledge AI Template, accessed February 22, 2026, <https://relevanceai.com/templates/retrieve-record-s-from-knowledge>
 44. Updating multiple Airtable records from JSON - Questions - Make Community, accessed February 22, 2026, <https://community.make.com/t/updating-multiple-airtable-records-from-json/51312>
 45. Map Array to HTTP "Make a Request" to Airtable API POST Call Not Working, accessed February 22, 2026, <https://community.make.com/t/map-array-to-http-make-a-request-to-airtable-api-post-call-not-working/46567>
 46. HOW to create multiple records from multiple outputs from API call into Airtable Base - Make Community, accessed February 22, 2026, <https://community.make.com/t/how-to-create-multiple-records-from-multiple-outputs-from-api-call-into-airtable-base/25266>
 47. Scenarios and Airtable Integration | Workflow Automation - Make, accessed February 22, 2026, <https://www.make.com/en/integrations/scenario-service/airtable>
 48. Build your AI Tools, accessed February 22, 2026, <https://relevanceai.com/tool>
 49. AWS re:Invent 2024 - Generative AI meets multi-tenancy: Inside a working solution (SAS407) - YouTube, accessed February 22, 2026,

<https://www.youtube.com/watch?v=hq3h5HNIBPE>

50. Personal access tokens - Airtable Web API, accessed February 22, 2026,
<https://airtable.com/developers/web/guides/personal-access-tokens>
51. Creating Personal Access Tokens - Airtable Support, accessed February 22, 2026,
<https://support.airtable.com/docs/creating-personal-access-tokens>
52. Connections - Help Center, accessed February 22, 2026,
<https://help.make.com/connections>
53. 6 Powerful Airtable Automation Tips Using Make.com - YouTube, accessed
February 22, 2026, <https://www.youtube.com/watch?v=phzwvyLrSCo>
54. Airtable batch update api call - Questions - Make Community, accessed February
22, 2026, <https://community.make.com/t/airtable-batch-update-api-call/27338>
55. What is the best ai answer API I should use for my N8N workflow - Reddit,
accessed February 22, 2026,
https://www.reddit.com/r/n8n/comments/1jkzgf5/what_is_the_best_ai_answer_api_i_should_use_for/